

Linux Kernel Self Protection Project


Linux Security Summit, Los Angeles
September 14, 2017

Kees (“Case”) Cook
keescook@chromium.org

<https://outflux.net/slides/2017/lss/kspp.pdf>

Agenda

- Background
 - “Security” in the context of this presentation
 - Why we need to change what we’re doing
 - Just fixing bugs isn’t sufficient
 - Upstream development model
- Kernel Self Protection Project
 - Who we are
 - What we’re doing
 - How you can help
- Challenges



Kernel Security

- More than access control (e.g. SELinux)
- More than attack surface reduction (e.g. seccomp)
- More than bug fixing (e.g. CVEs)
- More than protecting userspace
- More than kernel integrity
- This is about *Kernel Self Protection*



Devices using Linux

- Servers, laptops, cars, phones, ...
- >2,000,000,000 active Android devices in 2017
- Vast majority are running v3.4 (with v3.10 slowly catching up)
- Bug lifetimes are even longer than upstream
- “Not our problem”? None of this matters: even if upstream fixes every bug found, and the fixes are magically sent to devices, bug lifetimes are still huge.




Upstream Bug Lifetime


- In 2010 Jon Corbet researched security flaws, and found that the average time between introduction and fix was about 5 years.
- My analysis of Ubuntu CVE tracker for the kernel from 2011 through 2017:
 - Critical: 3 @ 5.3 years
 - High: 59 @ 6.4 years
 - Medium: 534 @ 5.6 years
 - Low: 273 @ 5.6 years



CVE lifetimes



critical & high CVE lifetimes



Upstream Bug Lifetime

- The risk is not theoretical. Attackers are watching commits, and they are better at finding bugs than we are:
 - <http://seclists.org/fulldisclosure/2010/Sep/268>
- Most attackers are not publicly boasting about when they found their 0-day...



Fighting Bugs

- We're finding them
 - Static checkers: compilers, coccinelle, sparse, smatch, coverity
 - Dynamic checkers: kernel, trinity, syzkaller, KASan-family
- We're fixing them
 - Ask Greg KH how many patches land in -stable
- They'll always be around
 - We keep writing them
 - They exist whether we're aware of them or not
 - Whack-a-mole is not a solution




Analogy: 1960s Car Industry

- @mricon's presentation at 2015 Linux Security Summit
 - <http://kernsec.org/files/lss2015/giant-bags-of-mostly-water.pdf>
- Cars were designed to run, not to fail
- Linux now where the car industry was in 1960s
 - <https://www.youtube.com/watch?v=fPF4fBGNK0U>
- We must handle failures (attacks) safely
 - Userspace is becoming difficult to attack
 - Containers paint a target on kernel
 - Lives depend on Linux



Killing bugs is nice

- Some truth to security bugs being “just normal bugs”
- Your security bug may not be my security bug
- We have little idea which bugs attackers use
- Bug might be in out-of-tree code
 - Un-upstreamed vendor drivers
 - Not an excuse to claim “not our problem”



Killing bug classes is better

- If we can stop an entire kind of bug from happening, we absolutely should do so!
- Those bugs never happen again
- Not even out-of-tree code can hit them
- But we'll never kill all bug classes



Killing exploitation is best

- We will always have bugs
- We must stop their exploitation
- Eliminate exploitation targets and methods
- Eliminate information leaks
- Eliminate anything that assists attackers
- *Even if it makes development more difficult*



Typical Exploit Chains

- Modern attacks tend to use more than one flaw
- Need to know where targets are
- Need to inject (or build) malicious code
- Need to locate malicious code
- Need to redirect execution to malicious code



What can we do?

- Many exploit mitigation technologies already exist (e.g. grsecurity/PaX) or have been researched (e.g. academic whitepapers), but many not present in upstream Linux kernel
- There is demand for kernel self-protection, and there is demand for it to exist in the upstream kernel
- <http://www.washingtonpost.com/sf/business/2015/11/05/net-of-in-security-the-kernel-of-the-argument/>

The Washington Post

Out-of-tree defenses?

- Some downstream kernel forks:
 - RedHat (ExecShield), Ubuntu (AppArmor), Android (Samsung KNOX), grsecurity (so many things)
- If you only *use* the kernel, and don't *develop* it, you're in a better position
 - But you're depending on a downstream fork
 - Fewer eyeballs (and less automated testing infrastructure) looking for vulnerabilities
 - Developing the kernel means using engineering resources for your fork
 - e.g. Android deals with multiple vendor forks already
 - Hard to integrate multiple forks
- Upstreaming means:
 - No more forward-porting
 - More review (never perfect, of course)



Digression 1: defending against email Spam

- Normal email server communication establishment:

Client	Server
[connect]	[accept]220 smtp.some.domain ESMTP ok
EHLO my.domain	250 ohai
MAIL FROM:<me@my.domain>	250 OK
RCPT TO:<you@your.domain>	250 OK
DATA	


Spam bot communication

- Success, and therefore timing, isn't important to Spam bots:

Client	Server
[connect]	[accept]220 smtp.some.domain ESMTP ok
EHLO my.domain	
MAIL FROM:<me@my.domain>	
RCPT TO:<you@your.domain>	
DATA	250 ohai
	250 OK
	250 OK

Trivially blocking Spam bots

- Insert a short starting delay



Powerful because it's not the default

- If everyone did this (i.e. it was upstream), bots would adapt
- If a defense is unexamined and/or only run by a subset of Linux users, it may be accidentally effective due to it being different, but may fail under closer examination
- Though, on the flip side, heterogeneous environments tend to be more resilient



Digression 2: Stack Clash research in 2017

- Underlying issues were identified in 2010
 - Fundamentally, if an attacker can control the memory layout of a setuid process, they may be able to manipulate it into colliding stack with other things, and arranging related overflows to gain execution control.
 - Linux tried to fix it with a 4K gap
 - grsecurity (from 2010 through at least their last public patch) took it further with a configurable gap, defaulting to 64K

A gap was not enough

- In addition to raising the gap size, grsecurity sensibly capped stack size of setuid processes, just in case:

```
do_execveat_common(...) {
    ...
    /* limit uid stack to 8MB
     * we saved the old limits above and will restore them if this exec fails */
    if (((!uid_eq(bprm->cred->euid, current_euid()) ||
        (!gid_eq(bprm->cred->egid, current_egid())))
        && (old_rlim[RLIMIT_STACK].rlim_cur > (8 * 1024 * 1024)))
        current->signal->rlim[RLIMIT_STACK].rlim_cur = 8 * 1024 * 1024;
    ...
}
```


Upstreaming the setuid stack size limit

- Landed in v4.14-rc1
- 15 patches
- Reviewed by at least 7 other people
- Made the kernel smaller
- Actually keeps the stack limited for setuid exec

```
16 files changed, 91 insertions(+), 159 deletions(-)
```

Important detail: threads

- Stack rlimit is a single value shared across entire thread-group
- Exec kills all other threads (part of the “point of no return”) as late in exec as possible
- If you check or set rlimits before the point of no return, you’re racing other threads



Un-upstreamed and unexamined for seven years

```
$ uname -r  
4.9.24-grsec+  
$ ulimit -s  
unlimited  
$ ls -la setuid-stack  
-rwsrwxr-x 1 root root 9112 Aug 11 09:17 setuid-stack  
$ ./setuid-stack  
Stack limit: 8388608  
$ ./raise-stack ./setuid-stack  
Stack limit: 18446744073709551615
```

Out-of-tree defenses need to be upstreamed

- While the preceding example isn't universally true for all out-of-tree defenses, it's a good example of why upstreaming is important, and why sometimes what looks like a tiny change turns into **much** more work.
- How do we get this done?



Kernel Self Protection Project

- <http://www.openwall.com/lists/kernel-hardening/>
 - <http://www.openwall.com/lists/kernel-hardening/2015/11/05/1>
- http://kernsec.org/wiki/index.php/Kernel_Self_Protection_Project
- People interested in coding, testing, documenting, and discussing the upstreaming of kernel self protection technologies and related topics.



Kernel Self Protection Project

- There are other people working on excellent technologies that ultimately revolve around the kernel protecting *userspace* from attack (e.g. brute force detection, SROP mitigations, etc)
- KSPP focuses on the kernel protecting the *kernel* from attack
- Currently ~12 organizations and ~10 individuals working on about ~20 technologies
- Slow and steady



Developers under KSPP umbrella

- LF's Core Infrastructure Initiative funded: Emese Revfy, with others pending
- Self-funded: Andy Lutomirski, Russell King, Valdis Kletnieks, Jason Cooper, Daniel Micay, David Windsor, Richard Weinberger, Richard Fellner, Daniel Gruss, Jason A. Donenfeld, Sandy Harris, Alexander Popov
- ARM: Catalin Marinas, Mark Rutland
- Canonical: Juerg Haefliger
- Cisco: Daniel Borkmann
- Docker: Tycho Andersen
- Google: Kees Cook, Thomas Garnier, Daniel Cashman, Jeff Vander Stoep, Jann Horn, Eric Biggers
- Huawei: Li Kun
- IBM: Michael Ellerman, Heiko Carstens, Christian Borntraeger
- Imagination Technologies: Matt Redfearn
- Intel: Elena Reshetova, Hans Liljestrand, Casey Schaufler, Michael Leibowitz, Dave Hansen, Peter Zijlstra
- Linaro: Ard Biesheuvel, David Brown, Arnd Bergmann
- Linux Foundation: Greg Kroah-Hartman
- Oracle: James Morris, Quentin Casasnovas, Yinghai Lu
- RedHat: Laura Abbott, Rik van Riel, Jessica Yu, Baoquan He

Probabilistic protections

- Protections that derive their strength from some system state being unknown to an attacker
- Weaker than “deterministic” protections since information exposures can defeat them, though they still have real-world value
- Familiar examples:
 - stack protector (canary value can be exposed)
 - Address Space Layout Randomization (offset can be exposed)



Deterministic protections

- Protections that derive their strength from organizational system state that always blocks attackers
- Familiar examples:
 - Read-only memory (writes will fail)
 - Bounds-checking (large accesses fail)



Bug class: stack overflow and exhaustion

Exploit example:

- <https://jon.oberheide.org/files/half-nelson.c>
- Mitigations:
 - **stack canaries**, e.g. gcc's **-fstack-protector (v2.6.30)** and **-fstack-protector-strong (v3.14)**
 - guard pages (e.g. GRKERNSEC_KSTACKOVERFLOW)
 - **vmap stack (v4.9 x86, v4.14 arm64)**, **removal of thread_info from stack (v4.9 x86, v4.10 arm64)**
 - *alloca checking* (e.g. **PAX_MEMORY_STACKLEAK**): *Alexander Popov*
 - shadow stacks (e.g. Clang SafeStack)

Bug class: integer over/underflow

- Exploit examples:
 - <https://cyseclabs.com/page?n=02012016>
 - <http://perception-point.io/2016/01/14/analysis-and-exploitation-of-a-linux-kernel-vulnerability-cve-2016-0728/>
- Mitigations:
 - check for atomic overflow (e.g. PAX_REFCOUNT)
 - *refcount_t: Elena Reshetova, David Windsor, Kees Cook, Ard Biesheuvel, Li Kun*
 - compiler plugin to detect multiplication overflows at runtime (e.g. PAX_SIZE_OVERFLOW)

Bug class: buffer overflows

- Exploit example:
 - <http://blog.includeseecurity.com/2014/06/exploit-walkthrough-cve-2014-0196-ptk-kernel-race-condition.html>
- Mitigations:
 - runtime validation of variable size vs copy_to_user / copy_from_user size (e.g. PAX_USERCOPY)
 - **CONFIG_HARDENED_USERCOPY (v4.8)**
 - *Usercopy whitelisting: David Windsor, Kees Cook*
 - *Usercopy slab segregation: David Windsor, Kees Cook*
 - metadata validation (e.g. glibc's heap protections)
 - **linked-list hardening (from grsecurity) CONFIG_DEBUG_LIST (v4.10)**
 - *CONFIG_SLUB_HARDENED, heap freelist obfuscation (from grsecurity): Daniel Micay, Kees Cook*
 - *Heap canaries: Daniel Micay*
 - FORTIFY_SOURCE (inspired by glibc), check buffer sizes of str*/mem* functions at compile- and run-time
 - **CONFIG_FORTIFY_SOURCE (v4.13)**
 - *Intra-object checking: Daniel Micay*

Bug class: format string injection

- Exploit example:
 - <http://www.openwall.com/lists/oss-security/2013/06/06/13>
- Mitigations:
 - **Drop %n entirely (v3.13)**
 - detect non-const format strings at compile time (e.g. gcc's -Wformat-security, or better plugin)
 - detect non-const format strings at run time (e.g. memory location checking done with glibc's -D_FORMAT_SOURCE=2)

Bug class: kernel pointer leak

- Exploit examples:
 - examples are legion: /proc (e.g. kallsyms, modules, slabinfo, iomem), /sys, **INET_DIAG (v4.1)**, etc
 - <http://vulnfactory.org/exploits/alpha-omega.c>
- Mitigations:
 - **kptr_restrict sysctl (v2.6.38)** too weak: requires dev opt-in
 - remove visibility to kernel symbols (e.g. GRKERNSEC_HIDESYM)
 - detect and block usage of %p or similar writes to seq_file or other user buffers (e.g. GRKERNSEC_HIDESYM + *PAX_USERCOPY*)

Bug class: uninitialized variables

- This is not just an information leak!
- Exploit example:
 - <https://outflux.net/slides/2011/defcon/kernel-exploitation.pdf>
- Mitigations:
 - *GCC plugin, stackleak: clear kernel stack between system calls (from PAX_MEMORY_STACKLEAK)*: Alexander Popov
 - **GCC plugin, structleak: instrument compiler to fully initialize all structures (from PAX_MEMORY_STRUCTLEAK)**: (`__user` v4.11, by-reference v4.14)

Bug class: use-after-free

- Exploit example:
 - <http://perception-point.io/2016/01/14/analysis-and-exploitation-of-a-linux-kernel-vulnerability-cve-2016-0728/>
- Mitigations:
 - clearing memory on free can stop attacks where there is no reallocation control (e.g. PAX_MEMORY_SANITIZE)
 - **Zero poisoning (v4.6)**
 - segregating memory used by the kernel and by userspace can stop attacks where this boundary is crossed (e.g. *PAX_USERCOPY*)
 - randomizing heap allocations can frustrate the reallocation efforts the attack needs to perform (e.g. OpenBSD malloc)
 - **Freelist randomization (SLAB: v4.7, SLUB: v4.8)**

Exploitation: finding the kernel

- Exploit examples (see “Kernel pointer leaks” above too):
 - <https://github.com/jonoberheide/ksymhunter>
- Mitigations:
 - hide symbols and kernel pointers (see “Kernel pointer leaks”)
 - kernel ASLR
 - text/modules base: **x86 (v3.14), arm64 (v4.6), MIPS (v4.7)**, ARM: *Ard Biesheuvel*
 - memory: **x86 (v4.8)**
 - PIE: **arm64 (v4.6)**, x86: *Thomas Garnier*
 - runtime randomization of kernel functions
 - executable-but-not-readable memory
 - **x86 (v4.6), arm64 (v4.9)**
 - per-build structure layout randomization (e.g. GRKERNSEC_RANDSTRUCT)
 - **manual (v4.13), automatic (v4.14)**

Exploitation: direct kernel overwrite

- How is this still a problem in the 21st century?
- Exploit examples:
 - Patch setuid to always succeed
 - <http://itszn.com/blog/?p=21> Overwrite vDSO
- Mitigations:
 - Executable memory cannot be writable (CONFIG_STRICT_KERNEL_RWX)
 - **s390: forever ago**
 - **x86: v3.18**
 - **ARM: v3.19**
 - **arm64: v4.0**

Exploitation: function pointer overwrite

- Also includes things like vector tables, descriptor tables (which can also be info leaks)
- Exploit examples:
 - <https://outflux.net/blog/archives/2010/10/19/cve-2010-2963-v4l-compat-exploit/>
 - https://blogs.oracle.com/ksplice/entry/anatomy_of_an_exploit_cve
- Mitigations:
 - read-only function tables (e.g. PAX_CONSTIFY_PLUGIN)
 - make sensitive targets that need one-time or occasional updates only writable during updates (e.g. PAX_KERNEXEC):
 - **__ro_after_init (v4.6)**

Exploitation: userspace execution

- Exploit example:
 - See almost all previous examples
- Mitigations:
 - hardware segmentation: **SMEP (x86)**, **PXN (ARM, arm64)**
 - emulated memory segmentation via page table swap, PCID, etc (e.g. **PAX_MEMORY_UDEREF**):
 - **Domains (ARM: v4.3)**
 - **TTBR0 (arm64: v4.10)**
 - *PCID (x86): Andy Lutomirski*
 - compiler instrumentation to set high bit on function calls

Exploitation: userspace data

- Exploit examples:
 - <https://github.com/geekben/towelroot/blob/master/towelroot.c>
 - <http://labs.bromium.com/2015/02/02/exploiting-badiret-vulnerability-cve-2014-9322-linux-kernel-privilege-escalation/>
- Mitigations:
 - hardware segmentation: **SMAP (x86), PAN (ARM, arm64)**
 - emulated memory segmentation via page table swap, PCID, etc (e.g. PAX_MEMORY_UDEREF):
 - **Domains (ARM: v4.3)**
 - **TTBR0 (arm64: v4.10)**
 - *PCID (x86): Andy Lutomirski*
 - *eXclusive Page Frame Ownership: Tycho Andersen, Juerg Haefliger*

Exploitation: reused code chunks

- Also known as Return Oriented Programming (ROP), Jump Oriented Programming (JOP), etc
- Exploit example:
 - <http://vulnfactory.org/research/h2hc-remote.pdf>
- Mitigations:
 - JIT obfuscation (e.g. BPF_HARDEN):
 - **eBPF JIT hardening (v4.7)**
 - compiler instrumentation for Control Flow Integrity (CFI):
 - Clang CFI <https://clang.llvm.org/docs/ControlFlowIntegrity.html>
 - kCFI <https://github.com/kcfi/docs>
 - GCC plugin: Return Address Protection, Indirect Control Transfer Protection (e.g. RAP)
<https://pax.grsecurity.net/docs/PaXTeam-H2HC15-RAP-RIP-ROP.pdf>

Added in v4.10

- PAN emulation, arm64
- `thread_info` relocated off stack, arm64
- Linked list hardening
- RNG seeding from UEFI, arm64
- W^X detection, arm64

Added in v4.11

- refcount_t infrastructure
- read-only usermodehelper
- structleak plugin

Added in v4.12

- read-only and fixed-location GDT, x86
- usercopy consolidation
- read-only LSM structures
- KASLR enabled by default, x86
- stack canary expanded to bit-width of host
- stack/heap gap expanded

Added in v4.13

- CONFIG_REFCOUNT_FULL
- CONFIG_FORTIFY_SOURCE
- randstruct (manual mode)
- ELF_ET_DYN_BASE lowered

Challenge: Culture

- Conservatism
 - 16 years to accept symlink restrictions upstream
- Responsibility
 - Kernel developers must accept the need for these changes
- Sacrifice
 - Kernel developers must accept the technical burden
- Patience
 - Out-of-tree developers must understand how kernel is developed

Challenge: Technical

- Complexity
 - Very few people are proficient at developing (much less debugging) these features
- Innovation
 - We must adapt the many existing solutions
 - We must create new technologies
- Collaboration
 - Explain rationale for new technologies
 - Make code understandable/maintainable by other developers and accessible across architectures

Challenge: Resources

- People
 - Dedicated developers
- People
 - Dedicated testers
- People
 - Dedicated backporters

A fluffy orange cat with white paws and chest, wearing a black leather collar with sharp silver spikes, is looking upwards and to the right. The background is plain white.

Thoughts?

Kees (“Case”) Cook

keescook@chromium.org

keescook@google.com

kees@outflux.net

<https://outflux.net/slides/2017/lss/kspp.pdf>

<http://www.openwall.com/lists/kernel-hardening/>

http://kernsec.org/wiki/index.php/Kernel_Self_Protection_Project